

```

1 import numpy as np
2 np.random.seed(123)
3 import pandas as pd
4 import pandas.testing as tm
5 import pandas.util.testing as tm
6 import seaborn as sns
7 import warnings
8 warnings.filterwarnings("ignore")
9 from tqdm import tqdm_notebook
10 from statsmodels.stats.outliers_influence import variance_inflation_factor
11 from sklearn.model_selection import train_test_split
12 from sklearn.decomposition import TruncatedSVD
13 from sklearn.decomposition import PCA
14 from sklearn.decomposition import FastICA
15 from sklearn.ensemble import RandomForestRegressor
16 from sklearn.model_selection import RandomizedSearchCV
17 import matplotlib.pyplot as plt
18 from sklearn.metrics import r2_score
19 import xgboost as xgb
20 from sklearn.tree import DecisionTreeRegressor
21 from sklearn.linear_model import LinearRegression
22 import statsmodels.api as sm
23 from sklearn.model_selection import cross_val_score
24 from sklearn.model_selection import GridSearchCV
25 from sklearn.linear_model import Ridge
26 from sklearn.linear_model import Lasso
27 from sklearn import linear_model
28 from sklearn import metrics
29 from sklearn.svm import SVR
30 import json
31 from tqdm import tqdm_notebook

```

 /usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:5: FutureWarning: pandas
 """

- Loading the DataSets

```

1 # loading the train and test dataset
2 df_1=pd.read_csv('/content/drive/My Drive/Colab Notebooks/22 Case Study 1/mercedes-benz

1 # converting categorical features to numerical features
2 x_dummies=pd.get_dummies(df_1, prefix_sep='_', drop_first=True)
3 x_dummies.head()
4 #https://pandas.pydata.org/pandas-docs/version/0.21.1/generated/pandas.get_dummies.html
5 #https://towardsdatascience.com/encoding-categorical-features-21a2651a065c

```



	ID	y	X10	X11	X12	X13	X14	X15	X16	X17	X18	X19	X20	X21	X22	X23
0	0	130.81	0	0	0	1	0	0	0	0	1	0	0	1	0	0
1	6	88.53	0	0	0	0	0	0	0	0	1	0	0	0	0	0
2	7	76.26	0	0	0	0	0	0	0	1	0	0	0	0	0	0

```

1 # creating a list of columns which have only zeros
2 zeros=[]
3 for i,j in x_dummies.any().items():#https://pandas.pydata.org/pandas-docs/stable/refere
4     if j==False:
5         zeros.append(i)

```

```

1 zeros # 'X339' is missing in my code
2 # we need to drop these columns

```

```

[> ['X11',
    'X93',
    'X107',
    'X233',
    'X235',
    'X268',
    'X289',
    'X290',
    'X293',
    'X297',
    'X330',
    'X347']]

```

```

1 #https://stackoverflow.com/questions/29294983/how-to-calculate-correlation-between-all-
2 # to delete columns with high correlation

```

- Dropping columns with only zeros

```

1 x_dummies = x_dummies.drop(zeros, axis=1)

```

- Dropping outliers => considering only IQR

```

1 x_filtered= x_dummies[x_dummies ['y']>70]#https://www.geeksforgeeks.org/drop-rows-from-
2 x_filtered= x_filtered[x_filtered['y']<150]
3 x_filtered.describe()

```

```

[>

```

	ID	y	X10	X12	X13	X14	
count	4194.000000	4194.000000	4194.000000	4194.000000	4194.000000	4194.000000	41
mean	4209.773724	100.439938	0.013352	0.074392	0.057940	0.428231	
std	2437.897217	11.994753	0.114792	0.262439	0.233658	0.494881	
min	0.000000	72.110000	0.000000	0.000000	0.000000	0.000000	

```

1 print(x_filtered.shape)
2 y=x_filtered['y']
3 x_filtered=x_filtered.drop(['y'], axis=1)
4 print(x_filtered.shape, y.shape)
5 print(type(x_filtered), type(y))

```

```

(4194, 545)
(4194, 544) (4194,)
<class 'pandas.core.frame.DataFrame'> <class 'pandas.core.series.Series'>

```

```

1 with open('/content/drive/My Drive/Colab Notebooks/22 Case Study 1/inf_indices.txt', 'r'
2     checc=json.loads(f.read())
3

```

```

1 inf_indices=checc

```

```

1 print(x_filtered.shape)
2 x_filtered = x_filtered.drop(inf_indices, axis=1)
3 #x_dummies_2= x_dummies_2.drop(inf_indices, axis=1)
4 print(x_filtered.shape)
5 #print(x_dummies_2.shape)

```

```

(4194, 544)
(4194, 188)

```

- Splitting data into Train and Test

```

1 from sklearn.model_selection import train_test_split
2 x_train, x_test, y_train, y_test = train_test_split(x_filtered, y, test_size=0.2, random_state=42)

```

<https://towardsdatascience.com/feature-extraction-techniques-d619b56e31be>

- From the above link we have learned that apart from SVD we can also add features using PCA, ICA, LDA, LLE
- Adding SVD features

```

1 from sklearn.decomposition import TruncatedSVD
2 tsvd= TruncatedSVD(n_components=2, random_state=42)
3 tsvd_train= tsvd.fit_transform(x_train)
4 tsvd_test= tsvd.transform(x_test)

```

```
1 tsvd_train.shape
```

```
(3355, 2)
```

- PCA

```
1 from sklearn.decomposition import PCA
2 pca = PCA(n_components=2, random_state=42)
3 pca_train= pca.fit_transform(x_train)
4 pca_test= pca.transform(x_test)
```

```
1 pca_train.shape
```

```
(3355, 2)
```

- ICA

```
1 from sklearn.decomposition import FastICA
2 ica=FastICA(n_components=2, random_state=42)
3 ica_train= ica.fit_transform(x_train)
4 ica_test= ica.transform(x_test)
```

```
1 ica_train.shape
```

```
(3355, 2)
```

- adding all these new features to the dataframe

```
1 for i in range(0, tsvd_train.shape[1]):
2     x_train['tsvd_'+str(i)]= tsvd_train[:, i]
3     x_test['tsvd_'+str(i)]= tsvd_test[:, i]
4     x_train['pca_'+str(i)]= pca_train[:, i]
5     x_test['pca_'+str(i)]= pca_test[:, i]
6     x_train['ica_'+str(i)]= ica_train[:, i]
7     x_test['ica_'+str(i)]= ica_test[:, i]
```

- Two way and Three way Feature interaction

```
1 x_train['X64 + X218']=x_train['X64']+x_train['X218']
2 x_test['X64 + X218']=x_test['X64']+x_test['X218']
3
4 x_train['X218 + X224 + X273']=x_train['X218']+x_train['X224'] + x_train['X273']
5 x_test['X218 + X224 + X273']=x_test['X218']+x_test['X224'] + x_test['X273']
6
7 x_train['X64 + X224 + X273']=x_train['X64']+x_train['X224'] + x_train['X273']
8 x_test['X64 + X224 + X273']=x_test['X64']+x_test['X224'] + x_test['X273']
9
```

```

10
11 x_train['X314 + X315']=x_train['X314']+x_train['X315']
12 x_test['X314 + X315']=x_test['X314']+x_test['X315']
13
14 x_train['X314 + X315 + X29']=x_train['X314']+x_train['X315'] + x_train['X29']
15 x_test['X314 + X315 + X29']=x_test['X314']+x_test['X315'] + x_test['X29']
16
17 x_train['X314 + X315 + X118']=x_train['X314']+x_train['X315'] + x_train['X118']
18 x_test['X314 + X315 + X118']=x_test['X314']+x_test['X315'] + x_test['X118']
19
20 x_train['X127 + X189']=x_train['X127']+x_train['X189']
21 x_test['X127 + X189']=x_test['X127']+x_test['X189']
22
23 x_train['X118 + X54']=x_train['X118']+x_train['X54']
24 x_test['X118 + X54']=x_test['X118']+x_test['X54']
25

1 print(x_train.shape, x_test.shape)

(3355, 202) (839, 202)

```

▼ 1: RandomForestRegressor

```

1 from sklearn.ensemble import RandomForestRegressor
2 # n_estimators, max_features, max_depth, min_samples_split, min_samples_leaf, bootstrap
3 from sklearn.model_selection import RandomizedSearchCV
4
5 # Number of trees in a Random forest
6 n_estimators=[int(x) for x in np.linspace (start=100, stop=1000, num=10)]
7
8 # Number of features to consider at every split
9 max_features=['auto', 'sqrt']
10
11 # Maximun no of levels in a tree
12 max_depth=[int(x) for x in np.linspace(10, 110, num = 11)]
13 max_depth.append(None)
14
15 # minimum number of samples required to split a node
16 min_samples_split = np.arange(50,250,20)
17
18 # Minimum number of samples required at each leaf node
19 min_samples_leaf = np.arange(5,50,5)
20 # Method of selecting samples for training each tree
21 bootstrap = [True, False]
22
23 # Create the random grid
24 random_grid = {'n_estimators': n_estimators,
25                 'max_features': max_features,
26                 'max_depth': max_depth,
27                 'min_samples_split': min_samples_split,
28                 'min_samples_leaf': min_samples_leaf,

```

```

29         'bootstrap': bootstrap}
30 print(random_grid)

```

```

↳ {'n_estimators': [100, 200, 300, 400, 500, 600, 700, 800, 900, 1000], 'max_features'

```

```

1 # Use the random grid to search for best hyperparameters
2 # First create the base model to tune
3 rf = RandomForestRegressor()
4 # Random search of parameters, using 3 fold cross validation,
5 # search across 100 different combinations, and use all available cores
6 rf_random = RandomizedSearchCV(estimator = rf, param_distributions = random_grid, n_ite
7 # Fit the random search model
8 rf_random.fit(x_train, y_train)

```

```

↳ Fitting 3 folds for each of 100 candidates, totalling 300 fits
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=-1)]: Done 37 tasks      | elapsed: 3.0min
[Parallel(n_jobs=-1)]: Done 158 tasks     | elapsed: 14.4min
[Parallel(n_jobs=-1)]: Done 300 out of 300 | elapsed: 28.3min finished
RandomizedSearchCV(cv=3, error_score=nan,
                    estimator=RandomForestRegressor(bootstrap=True,
                                                        ccp_alpha=0.0,
                                                        criterion='mse',
                                                        max_depth=None,
                                                        max_features='auto',
                                                        max_leaf_nodes=None,
                                                        max_samples=None,
                                                        min_impurity_decrease=0.0,
                                                        min_impurity_split=None,
                                                        min_samples_leaf=1,
                                                        min_samples_split=2,
                                                        min_weight_fraction_leaf=0.0,
                                                        n_estimators=100,
                                                        n_jobs=None, oob_score=Fals...
                    'max_depth': [10, 20, 30, 40, 50, 60,
                                   70, 80, 90, 100, 110,
                                   None],
                    'max_features': ['auto', 'sqrt'],
                    'min_samples_leaf': array([ 5, 10, 15, 20, 25,
                    'min_samples_split': array([ 50, 70, 90, 110,
                    'n_estimators': [100, 200, 300, 400,
                                       500, 600, 700, 800,
                                       900, 1000]},
                    pre_dispatch='2*n_jobs', random_state=42, refit=True,
                    return_train_score=False, scoring=None, verbose=2)

```

```

1 rf_random.best_params_

```

```

↳ {'bootstrap': False,
    'max_depth': 60,
    'max_features': 'sqrt',
    'min_samples_leaf': 5,
    'min_samples_split': 150,
    'n_estimators': 200}

```

```

1 best_rf=rf_random.best_estimator_
2 best_rf

```


```
↳ RandomForestRegressor(bootstrap=False, ccp_alpha=0.0, criterion='mse',
                          max_depth=60, max_features='sqrt', max_leaf_nodes=None,
                          max_samples=None, min_impurity_decrease=0.0,
                          min_impurity_split=None, min_samples_leaf=5,
                          min_samples_split=150, min_weight_fraction_leaf=0.0,
                          n_estimators=200, n_jobs=None, oob_score=False,
                          random_state=None, verbose=0, warm_start=False)
```

```
1 #temp
2 best_rf= RandomForestRegressor(bootstrap=False, ccp_alpha=0.0, criterion='mse',
3                               max_depth=60, max_features='sqrt', max_leaf_nodes=None,
4                               max_samples=None, min_impurity_decrease=0.0,
5                               min_impurity_split=None, min_samples_leaf=5,
6                               min_samples_split=150, min_weight_fraction_leaf=0.0,
7                               n_estimators=200, n_jobs=None, oob_score=False,
8                               random_state=None, verbose=0, warm_start=False)
```

```
1 # fitting the best model on to the dataset
2 best_rf.fit(x_train, y_train)
3 y_train_pred=best_rf.predict(x_train)
4 y_test_pred=best_rf.predict(x_test)
```

```
1 features = x_train.columns
2 importances = best_rf.feature_importances_
3 indices = (np.argsort(importances))[-40:]
4 plt.figure(figsize=(10,12))
5 plt.title('Feature Importances')
6 plt.barh(range(len(indices)), importances[indices], color='r', align='center')
7 plt.yticks(range(len(indices)), [features[i] for i in indices])
8 plt.xlabel('Relative Importance')
9 plt.grid()
10 plt.show()
```

↳



```

1 # checking r2 score
2 from sklearn.metrics import r2_score
3 print('train r2_score', r2_score(y_train, y_train_pred))
4 print('='*30)
5 print('test r2_score', r2_score(y_test, y_test_pred))

```

```

train r2_score 0.6574681408333823
=====
test r2_score 0.6483743021527534

```

0.6574681408333823

2: XGBoostRegressor

```
1 import xgboost as xgb
```

```

1 xgb_model=xgb.XGBRegressor()
2 xgb_model.fit(x_train, y_train)
3 y_pred=xgb_model.predict(x_train)
4 learning_rate=[0.001, 0.01, 0.1, 0.2, 0.3]
5 n_estimators=[50, 100, 150, 200, 250, 300]
6 hyperparameters= dict(learning_rate=learning_rate, n_estimators=n_estimators)
7

```

```
[05:33:46] WARNING: /workspace/src/objective/regression_obj.cu:152: reg:linear is now
```

```

1 from sklearn.model_selection import RandomizedSearchCV
2 rsearch = RandomizedSearchCV(xgb_model, hyperparameters, n_iter=10, random_state=41)
3 rsearch.fit(x_train, y_train)

```

```
1 print(rsearch.best_params_)
```



```
➦ {'n_estimators': 150, 'learning_rate': 0.1}
1 from sklearn.metrics import r2_score
2 train_r2=[]
3 cv_r2=[]
4 dict_cv_r2={}
5 learning_rate=[0.3, 0.2, 0.1, 0.01, 0.001]
6 n_estimators=[50, 100, 150, 200, 250, 300]
7 for i in tqdm_notebook(learning_rate):
8     for j in tqdm_notebook(n_estimators):
9         xb=xgb.XGBRegressor(learning_rate=i,n_estimators=j)
10        xb.fit(x_train,y_train)
11        y_cap_train=xb.predict(x_train)
12        y_cap_cv=xb.predict(x_test)
13        train_r2.append(r2_score(y_train,y_cap_train))
14        k=r2_score(y_test,y_cap_cv)
15        dict_cv_r2[k]=i,j
16        cv_r2.append(r2_score(y_test,y_cap_cv))
17 maximum_auc_score=max(cv_r2)
18 print(maximum_auc_score)
```



```
HBox(children=(FloatProgress(value=0.0, max=5.0), HTML(value='')))
HBox(children=(FloatProgress(value=0.0, max=6.0), HTML(value='')))
[05:35:56] WARNING: /workspace/src/objective/regression_obj.cu:152: reg:linear is now
[05:35:57] WARNING: /workspace/src/objective/regression_obj.cu:152: reg:linear is now
[05:35:58] WARNING: /workspace/src/objective/regression_obj.cu:152: reg:linear is now
```

```
1 print(dict_cv_r2)
2 dict_cv_r2[maximum_auc_score]
```

```
↳ {0.6402537526049236: (0.3, 50), 0.6252262886435493: (0.3, 100), 0.6157867395052805: (
(0.1, 50)
```

```
[05:36:15] WARNING: /workspace/src/objective/regression_obj.cu:152: reg:linear is now
```

```
1 best_xgb=rsearch.best_estimator_
2 best_xgb
```

```
↳ XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
               colsample_bynode=1, colsample_bytree=1, gamma=0,
               importance_type='gain', learning_rate=0.1, max_delta_step=0,
               max_depth=3, min_child_weight=1, missing=None, n_estimators=150,
               n_jobs=1, nthread=None, objective='reg:linear', random_state=0,
               reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
               silent=None, subsample=1, verbosity=1)
```


```
1 # fitting the best model on to the dataset
2 best_xgb.fit(x_train, y_train)
3 y_train_pred=best_xgb.predict(x_train)
4 y_test_pred=best_xgb.predict(x_test)
```

```
↳ [05:37:21] WARNING: /workspace/src/objective/regression_obj.cu:152: reg:linear is now
```

```
HBox(children=(FloatProgress(value=0.0, max=6.0), HTML(value='')))
```

```
1 import matplotlib.pyplot as plt
2 features = x_train.columns
3 importances = best_xgb.feature_importances_
4 indices = (np.argsort(importances))[-40:]
5 plt.figure(figsize=(10,12))
6 plt.title('Feature Importances')
7 plt.barh(range(len(indices)), importances[indices], color='r', align='center')
8 plt.yticks(range(len(indices)), [features[i] for i in indices])
9 plt.xlabel('Relative Importance')
10 plt.grid()
11 plt.show()
```

```
↳
```



```

1 # checking r2 score
2 from sklearn.metrics import r2_score
3 print('train r2_score', r2_score(y_train, y_train_pred))
4 print('='*30)
5 print('test r2_score', r2_score(y_test, y_test_pred))

```

```

train r2_score 0.6931087582848501
=====
test r2_score 0.6510758493146287

```

3: DecisionTreeRegressor

```

1 from sklearn.tree import DecisionTreeRegressor
2 dt_model=DecisionTreeRegressor()
3 dt_model

```

```

DecisionTreeRegressor(ccp_alpha=0.0, criterion='mse', max_depth=None,
                      max_features=None, max_leaf_nodes=None,
                      min_impurity_decrease=0.0, min_impurity_split=None,
                      min_samples_leaf=1, min_samples_split=2,
                      min_weight_fraction_leaf=0.0, presort='deprecated',
                      random_state=None, splitter='best')

```

```

1 max_depth=[1, 5, 10, 50]
2 min_samples_split=[5, 10, 100, 500]
3 hyperparameters= dict(max_depth=max_depth, min_samples_split=min_samples_split)

```

```

3 hyperparameters = dict(max_depth=max_depth, min_samples_split=min_samples_split,
4

```

```

1 from sklearn.model_selection import RandomizedSearchCV
2 rsearch = RandomizedSearchCV(dt_model, hyperparameters, n_iter=10, random_state=41)
3 rsearch.fit(x_train, y_train)
4 print(rsearch.best_params_)

```

```

↳ {'min_samples_split': 500, 'max_depth': 5}

```

```

1 rsearch.best_params_

```

```

↳ {'max_depth': 5, 'min_samples_split': 500}

```

```

1 best_dt=rsearch.best_estimator_
2 best_dt

```

```

↳ DecisionTreeRegressor(ccp_alpha=0.0, criterion='mse', max_depth=5,
                        max_features=None, max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=500,
                        min_weight_fraction_leaf=0.0, presort='deprecated',
                        random_state=None, splitter='best')

```

```

1 # fitting the best model on to the dataset
2 best_dt.fit(x_train, y_train)
3 y_train_pred=best_dt.predict(x_train)
4 y_test_pred=best_dt.predict(x_test)

```

```

1 import matplotlib.pyplot as plt
2 features = x_train.columns
3 importances = best_dt.feature_importances_
4 indices = (np.argsort(importances))[-40:]
5 plt.figure(figsize=(10,12))
6 plt.title('Feature Importances')
7 plt.barh(range(len(indices)), importances[indices], color='r', align='center')
8 plt.yticks(range(len(indices)), [features[i] for i in indices])
9 plt.xlabel('Relative Importance')
10 plt.grid()
11 plt.show()

```

```

↳

```



```

1 # checking r2 score
2 from sklearn.metrics import r2_score
3 print('train r2_score', r2_score(y_train, y_train_pred))
4 print('='*30)
5 print('test r2_score', r2_score(y_test, y_test_pred))

```

```

train r2_score 0.6290528360587266
=====
test r2_score 0.6554069343566447

```

0.0 0.1 0.2 0.3 0.4 0.5 0.6 0.7 0.8

4: Ridge Regressor

```

1 from sklearn.linear_model import Ridge
2 ridge_regressor= Ridge()
3 ridge_regressor

```

```

Ridge(alpha=1.0, copy_X=True, fit_intercept=True, max_iter=None,
      normalize=False, random_state=None, solver='auto', tol=0.001)

```

```

1 parameters= {'alpha': [10,20, 50, 100, 150, 200, 250, 300]}
2 ridge_grid_scv= GridSearchCV(ridge_regressor, parameters, scoring='r2', cv=5)
3 ridge_grid_scv.fit(x_train, y_train)

```

```


```

```
GridSearchCV(cv=5, error_score=nan,
             estimator=Ridge(alpha=1.0, copy_X=True, fit_intercept=True,
                             max_iter=None, normalize=False, random_state=None,
                             solver='auto', tol=0.001),
             iid='deprecated', n_jobs=None,
             param_grid={'alpha': [10, 20, 50, 100, 150, 200, 250, 300]}),
```

```
1 print(ridge_grid_scv.best_params_)
2 print(ridge_grid_scv.best_score_)
```

```
↳ {'alpha': 50}
0.5954015864151923
```

```
1 ridge_regressor= Ridge(alpha=50)
2 ridge_regressor
```

```
↳ Ridge(alpha=50, copy_X=True, fit_intercept=True, max_iter=None, normalize=False,
        random_state=None, solver='auto', tol=0.001)
```

```
1 ridge_regressor.fit(x_train, y_train)
2 y_train_pred=ridge_regressor.predict(x_train)
3 y_test_pred =ridge_regressor.predict(x_test)
```

```
1 # checking r2 score
2 from sklearn.metrics import r2_score
3 print('train r2_score', r2_score(y_train, y_train_pred))
4 print('='*30)
5 print('test r2_score', r2_score(y_test, y_test_pred))
```

```
↳ train r2_score 0.622606917016735
=====
test r2_score 0.6398563894858129
```

▼ 5: Lasso Regressor

```
1 lasso= Lasso()
2 lasso
```

```
↳ Lasso(alpha=1.0, copy_X=True, fit_intercept=True, max_iter=1000,
        normalize=False, positive=False, precompute=False, random_state=None,
        selection='cyclic', tol=0.0001, warm_start=False)
```

```
1 parameters= {'alpha': [1e-15, 1e-10, 1e-8, 1e-4, 1e-3, 1e-2, 1, 5, 10, 20]}
2 lasso_grid_scv= GridSearchCV(lasso, parameters, scoring='r2', cv=5)
3 lasso_grid_scv.fit(x_train, y_train)
```

```
↳
```

```

GridSearchCV(cv=5, error_score=nan,
              estimator=Lasso(alpha=1.0, copy_X=True, fit_intercept=True,
                              max_iter=1000, normalize=False, positive=False,
                              precompute=False, random_state=None,
                              selection='cyclic', tol=0.0001, warm_start=False),

1 print(lasso_grid_scv.best_params_)
2 print(lasso_grid_scv.best_score_)

↳ {'alpha': 0.01}
   0.596278101998697

1 lasso_regressor= Lasso(alpha=0.01)
2 lasso_regressor

↳ Lasso(alpha=0.01, copy_X=True, fit_intercept=True, max_iter=1000,
        normalize=False, positive=False, precompute=False, random_state=None,
        selection='cyclic', tol=0.0001, warm_start=False)

1 #https://towardsdatascience.com/how-to-perform-lasso-and-ridge-regression-in-python-3b3
2 #https://www.analyticsvidhya.com/blog/2016/01/ridge-lasso-regression-python-complete-tu
3 #https://alfurka.github.io/2018-11-18-grid-search/
4 lasso_regressor.fit(x_train, y_train)
5 y_train_pred=lasso_regressor.predict(x_train)
6 y_test_pred =lasso_regressor.predict(x_test)

1 # checking r2 score
2 from sklearn.metrics import r2_score
3 print('train r2_score', r2_score(y_train, y_train_pred))
4 print('='*30)
5 print('test r2_score', r2_score(y_test, y_test_pred))

↳ train r2_score 0.6256643013649388
   =====
   test r2_score 0.6463713208948576

```

▼ 6: AdaBoostRegressor

```

1 from sklearn.ensemble import AdaBoostRegressor
2 from sklearn.svm import SVR
3 abr= AdaBoostRegressor()
4 abr

1 parameters = {'learning_rate':[0.0001, 0.001, 0.01, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7,
2 gsearch = GridSearchCV(abr, parameters, scoring='r2')
3 gsearch.fit(x_train, y_train)

↳

```



```

GridSearchCV(cv=None, error_score=nan,
             estimator=AdaBoostRegressor(base_estimator=None, learning_rate=1.0,
                                         loss='linear', n_estimators=50,
                                         random_state=None),
             iid='deprecated', n_jobs=None,
             param_grid={'learning_rate': [0.0001, 0.001, 0.01, 0.1, 0.2, 0.3,
                                           0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1]},

```

```
1 print(gsearch.best_estimator_)
```

```
2 print(gsearch.best_score_)
```

```

↳ AdaBoostRegressor(base_estimator=None, learning_rate=0.0001, loss='exponential',
                    n_estimators=200, random_state=None)
0.6106137543959904

```

```

1 abr= AdaBoostRegressor(base_estimator=None, learning_rate=0.0001, loss='exponential',
2                       n_estimators=200, random_state=None)
3 abr

```

```

↳ AdaBoostRegressor(base_estimator=None, learning_rate=0.0001, loss='exponential',
                    n_estimators=200, random_state=None)

```

```

1 abr.fit(x_train, y_train)
2 y_train_pred=abr.predict(x_train)
3 y_test_pred =abr.predict(x_test)

```

```

1 # checking r2 score
2 from sklearn.metrics import r2_score
3 print('train r2_score', r2_score(y_train, y_train_pred))
4 print('='*30)
5 print('test r2_score', r2_score(y_test, y_test_pred))

```

```

↳ train r2_score 0.6134014860036547
=====
test r2_score 0.653160564280501

```